

CMSE 381, Fundamental Data Science Methods

September 7, 2025

Homework 1

Lowell Monis

Question 1: ISLP § 2.4.9

```
[2]: auto = pd.read_csv('../data/Auto.csv')
auto=auto.replace('?', np.nan)
auto=auto.dropna()
auto['horsepower']=auto['horsepower'].astype('int')
auto=auto.reset_index(drop=True)
```

(a) Which of the predictors are quantitative, and which are qualitative? The following predictors are qualitative: **name** and **origin**.

The following predictors are quantitative: **mpg**, **cylinders**, **displacement**, **horsepower**, **weight**, **acceleration**, and **year**.

Why do I think so? For the quantitative predictors, the reasoning is rather simpler. All of them have multiple unique values and are measurable quantities. Even if they are fewer in number, say below 100, they are not really placed equidistantly to be considered discrete. They also have much variation. In the case of **year**, even if it has 13 unique values and they are equidistant from each other, it is a timescale, and timescale values are usually considered quantitative since they do not qualitatively or categorically describe the data in most cases, and time is also a measurable quantity. In the case of **cylinder**, even with 5 unique values, it provides a count of the number of cylinders, thus showing an ability to be measured. There is a limit to the number of cylinders, and cylinders cannot be partial, so even with limited unique values and the appearance of discreteness, it quantitatively describes the data.

Now, even though **name** has the most unique values, being an indication of high variation in the data, it is a text-based variable, thus describing the data qualitatively. It is the name of the automobile, which cannot be used to quantitatively describe the data, and cannot be measured as a quantity. **origin** is the one variable that has the least variation and is discrete, like the perfect categorical variable for clustering or classification. But the primary reasoning for classifying this as qualitative is the fact that it describes where the car was made in three levels. A choice was made to represent these countries using numbers rather than the country name. Origin, like the name, is not a measurable quantity.

```
[3]: for i in auto.columns:
      print("There are", len(auto[i].unique()), "unique entries in predictor", i)
```

```
There are 127 unique entries in predictor mpg
There are 5 unique entries in predictor cylinders
There are 81 unique entries in predictor displacement
There are 93 unique entries in predictor horsepower
There are 346 unique entries in predictor weight
There are 95 unique entries in predictor acceleration
There are 13 unique entries in predictor year
There are 3 unique entries in predictor origin
There are 301 unique entries in predictor name
```

(b) What is the range of each quantitative predictor? You can answer this using the `min()` and `max()` methods in `numpy`. I used the `describe()` method and selected specific rows rather than repetitively using a command. The range can be read from top to bottom (minimum to maximum).

```
[4]: auto.drop('origin', axis=1).describe().loc[['min', 'max']]
```

```
[4]:      mpg  cylinders  displacement  horsepower  weight  acceleration  year
min    9.0         3.0          68.0         46.0  1613.0           8.0  70.0
max   46.6         8.0         455.0        230.0  5140.0          24.8  82.0
```

(c) What is the mean and standard deviation of each quantitative predictor? I used the same technique here as I did in (b).

```
[5]: auto.drop('origin', axis=1).describe().loc[['mean', 'std']]
```

```
[5]:      mpg  cylinders  displacement  horsepower  weight \
mean  23.445918   5.471939   194.411990   104.469388  2977.584184
std    7.805007   1.705783   104.644004    38.491160   849.402560

      acceleration  year
mean    15.541327  75.979592
std     2.758864   3.683737
```

(d) Now remove the 10th through 85th observations. What is the range, mean, and standard deviation of each predictor in the subset of the data that remains? The range, once again, is read from minimum to maximum values here.

```
[6]: auto.drop(range(9,85), axis=0).describe().loc[['min', 'max', 'mean', 'std']]
```

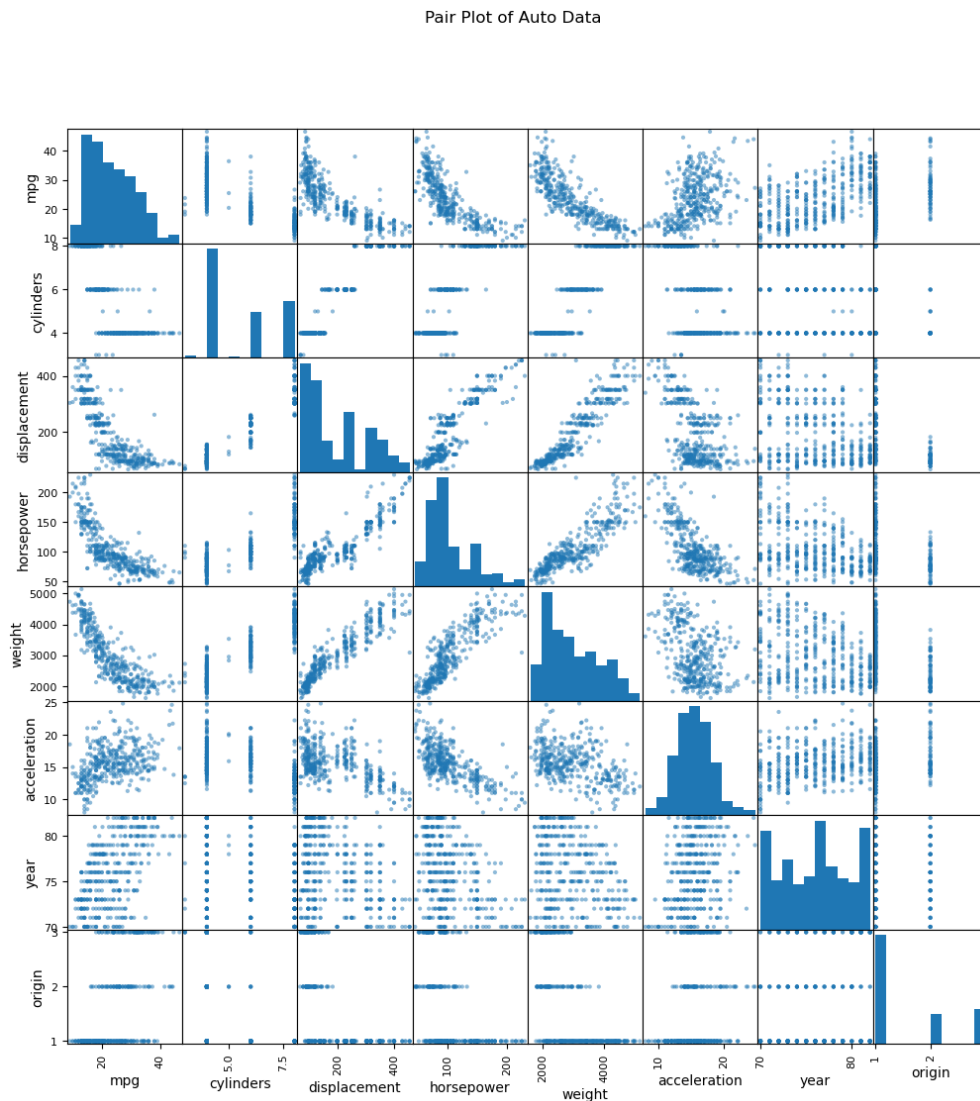
```
[6]:      mpg  cylinders  displacement  horsepower  weight \
min    11.000000   3.000000   68.000000   46.000000  1649.000000
max    46.600000   8.000000  455.000000  230.000000  4997.000000
mean    24.404430   5.373418  187.240506  100.721519  2935.971519
std     7.867283   1.654179   99.678367   35.708853   811.300208
```

	acceleration	year	origin
min	8.500000	70.000000	1.000000
max	24.800000	82.000000	3.000000
mean	15.726899	77.145570	1.601266
std	2.693721	3.106217	0.819910

(e) Using the full data set, investigate the predictors graphically, using scatterplots or other tools of your choice. Create some plots highlighting the relationships among the predictors. Comment on your findings.

```
[7]: pd.plotting.scatter_matrix(auto, figsize=(12, 12))
plt.suptitle('Pair Plot of Auto Data')
```

```
[7]: Text(0.5, 0.98, 'Pair Plot of Auto Data')
```



The pair plot provides a comprehensive visual summary of the relationships between all the variables. The scatter plots reveal several key findings:

- There's a strong positive correlation between several of the quantitative predictors. For example, **horsepower**, **weight**, and **displacement** are all positively related to each other. Cars that have larger engine displacement tend to be heavier and have more horsepower. The relationship appears mostly linear.
- The **cylinders** variable, which is discrete, shows a clear positive relationship with other variables like **displacement**, **horsepower**, and **weight**. Vehicles with more cylinders generally have larger, heavier, and more powerful engines.
- The **acceleration** variable has negative trends with **displacement** (slightly), **horsepower** (fairly), and **weight** (slightly), suggesting that heavier, more powerful cars may not always be the fastest.

(f) Suppose that we wish to predict gas mileage (mpg) on the basis of the other variables. Do your plots suggest that any of the other variables might be useful in predicting mpg? Justify your answer. The plots strongly suggest that many of the other variables will be useful in predicting mpg. We can justify this by observing the relationships in the first column of the pair plot, which shows mpg plotted against every other variable.

- The scatter plots of mpg versus **horsepower**, **weight**, and **displacement** show a strong, clear negative relationship. As **horsepower**, **weight**, or **displacement** increase, mpg consistently decreases. The distinct downward slope in these plots indicates that they are highly useful predictors.
- The mpg vs. **cylinders** plot shows that the average mpg value is significantly lower for cars with more cylinders compared to those with fewer.

In short, any variable that shows a discernible trend (linear or otherwise) with mpg is a useful predictor, and the pair plot confirms that most of the variables in this dataset fit that description.

Question 2: ISLP § 2.4.2

Explain whether each scenario is a classification or regression problem, and indicate whether we are most interested in inference or prediction. Finally, provide n and p . In each of these scenarios, we establish that n is the number of data points, and p is the number of predictors, or features.

(a) We collect a set of data on the top 500 firms in the US. For each firm we record profit, number of employees, industry and the CEO salary. We are interested in understanding which factors affect CEO salary.

In this scenario, an attempt is being made to understand the behaviors of the data and its effect on the target variable within the time frame provided, rather than predict behaviors of the target variable in the future. I can thus conclude that we are interested primarily in **inference** here. Considering that salary is a quantitative variable, due to its variability, wide range, and continuous, rather than discrete, numerical nature as a value, I can assume that this analysis would need **regression**, since classification will need a few discrete, and a certain number of groups to classify into. Here, $n = 500$, since the analysis is using data from 500 firms, and $p = 3$, considering the CEO salary is a target variable and won't be counted as a predictor.

(b) We are considering launching a new product and wish to know whether it will be a success or a failure. We collect data on 20 similar products that were previously launched. For each product we have recorded whether it was a success or failure, price charged for the product, marketing budget, competition price, and ten other variables.

In this scenario, an attempt is being made to forecast the outcome of a new product through different market factors of existing products that are similar to the new one. I can thus conclude that we are interested primarily in **prediction** here, rather than understanding what led to the success and failure of the 20 other products that have already been launched. Considering that the target variable, which is the market outcome of the product, has two discrete levels, i.e., success or failure, we assume this variable is of a qualitative nature, and we can further assume that this is a **classification** problem, since we have to use the other factors to classify the new product between one of two levels. Here, $n = 20$, since there are twenty existing data points. $p = 13$ here; this is because there are 4 named variables, and 10 unnamed ones, out of which 1 of the variables is the target variable.

(c) We are interested in predicting the % change in the USD/Euro exchange rate in relation to the weekly changes in the world stock markets. Hence we collect weekly data for all of 2012. For each week we record the % change in the USD/Euro, the % change in the US market, the % change in the British market, and the % change in the German market.

In this scenario, an attempt is being made to forecast the percent change in the exchange rate in relation to the percent change by week in three other markets around the world. I can thus conclude, from just the framing of the question, that this is a **prediction** problem. Considering that percent changes are on a continuous numerical scale and fluctuate between multiple decimal values, this is a quantitative target variable, and thus leads me to conclude that this problem is a **regression** problem. There are weekly values for the year 2012 in the data here, so $n = 52$, since there were 52 weeks in the year 2012. Considering we are using market changes from three markets to predict the exchange rate, we can also conclude that $p = 3$.

Question 3: ISLP § 2.4.4

(a) Describe *one* real-life application in which classification might be useful. Describe the response, as well as the predictors. Is the goal of each application inference or prediction? Explain your answer. Here are three applications of classification I see in real life:

1. While this is a massive oversimplification, the “algorithm” in social media and video-sharing websites can be thought of as a classification model. The response here is a discrete, categorical variable with levels “recommend” and “do not recommend”. The predictors could be things like what videos I have watched, for how long, whether I finished them, and if I have watched them again (these apply to video-sharing sites); what posts I have liked, commented on, or shared; the accounts I follow/am subscribed to; my demographic information, like age, gender, region, etc. This is clearly a prediction, by virtue of the fact that the model is predicting what I want to watch or see, as well as the fact that the model can be entirely wrong. An example is when I am recommended posts not in line with my political views, just because I may have searched an account that promotes views to inform myself.
2. I believe one of the most common applications of classification is the spam filter in emails. The response here is a discrete, categorical variable with levels “spam” and “not spam”. The

predictors could be things like the occurrence of some common words associated with a spam email (e.g., “free,” “winner,” “urgent”), the sender’s address (this might need NLP, however, if my understanding of NLP is correct), and the number of recipients. This is a prediction, by virtue of the fact that the goal is to automatically sort an email before the user even sees it. The model’s primary task is to make a fast and accurate guess on new emails, and it can be entirely wrong, sometimes sending a legitimate email to the spam folder.


3. Another important application is in loan or credit card applications to analyze risk in lending money to the borrower or approving an applicant for their credit card. The response here is a discrete, categorical variable with levels “risky” and “not risky”, or “approve” and “deny”. The predictors can include credit score, income, past defaults, age, and employment status. While I feel this is obviously a prediction model, a perspective can be provided in terms of inference by virtue of the fact that the model can be used to rewrite lending policies based on past data.

I used Gemini by Google for the first application, since I was a little lost and needed a head start on what were potential predictors for a model like this. My prompt was as follows: > what are examples of predictors for the “algorithm” that recommends content on social media and video-sharing sites, in the context that i am simplifying this model to a simple classification problem?

Note to grader: I just read the question properly, but I feel bad deleting the other responses I wrote, so I am submitting three examples for the first one.

(b) Describe *one* real-life application in which regression might be useful. Describe the response, as well as the predictors. Is the goal of each application inference or prediction? Explain your answer. I believe fitness apps can use regression to see how their app helps their users’ fitness goals through responses such as a net change in weight, for example, where the user wants to undergo weight loss. Here, predictors can be the frequency of active time, number of workouts, intensity of workouts, original weight, goal weight, etc. This is an inference application, since the app can be tweaked based on the net percent change in weight by the user.

Question 4: ISLP § 2.4.3



(a) Provide a sketch of typical (squared) bias, variance, training error, test error, and irreducible error curves, on a single plot, as we go from less flexible statistical learning methods towards more flexible approaches. The x-axis should represent the amount of flexibility in the method, and the y-axis should represent the values for each curve. There should be five curves. Make sure to label each one.

(b) Explain why each of the five curves has the shape displayed in part (a). As we increase the flexibility of a statistical learning method, there's a trade-off between bias and variance. Less flexible models, such as regression-based models, have high bias because they can't capture the complex patterns in the data; however, they have low variance because their predictions remain consistent even with different training sets. Conversely, highly flexible models, and even black-box models, like deep neural networks, have low bias because they fit the training data extremely well, but they suffer from high variance as they become overly sensitive to noise. Thus, the bias and variance curves are almost opposite to each other, with bias decreasing with nearly, if not exactly, the same magnitude as variance increases.

The bias-variance trade-off impacts the error curves. The training error consistently decreases as flexibility increases because the model can fit the training data more closely. However, the test error, which is what we should care about, follows a near U-shaped curve. This is because it decreases as bias is reduced, but eventually rises again as the model becomes too flexible and suffers from high variance and overfitting. The lowest point on this curve represents the optimal balance. Finally, there is an intrinsic, constant, irreducible error that no model can ever eliminate, as it's caused by the random noise in the data itself, and not the noise that is captured by highly flexible, overfitted models.

Question 5: ISLP § 2.4.8

This exercise relates to the College data set, which can be found in the file `College.csv` on the book website. It contains a number of variables for 777 different universities and colleges in the US. The variables are

- **Private**: Public/private indicator
- **Apps**: Number of applications received
- **Accept**: Number of applicants accepted
- **Enroll**: Number of new students enrolled
- **Top10perc**: New students from top 10% of high school class
- **Top25perc**: New students from top 25% of high school class
- **F.Undergrad**: Number of full-time undergraduates
- **P.Undergrad**: Number of part-time undergraduates
- **Outstate**: Out-of-state tuition
- **Room.Board**: Room and board costs
- **Books**: Estimated book costs
- **Personal**: Estimated personal spending
- **PhD**: Percent of faculty with Ph.D.s
- **Terminal**: Percent of faculty with terminal degree
- **S.F.Ratio**: Student/faculty ratio
- **perc.alumni**: Percent of alumni who donate
- **Expend**: Instructional expenditure per student
- **Grad.Rate**: Graduation rate

Before reading the data into Python, it can be viewed in Excel or a text editor.

(a) Use the `pd.read_csv()` function to read the data into Python. Call the loaded data `college`. Make sure that you have the directory set to the correct location for the data.

```
[8]: college = pd.read_csv('../data/College.csv')
```

(b) Look at the data used in the notebook by creating and running a new cell with just the code `college` in it.

```
[9]: college
```

```
[9]:
```

	Unnamed: 0	Private	Apps	Accept	Enroll	Top10perc	\
0	Abilene Christian University	Yes	1660	1232	721	23	
1	Adelphi University	Yes	2186	1924	512	16	
2	Adrian College	Yes	1428	1097	336	22	
3	Agnes Scott College	Yes	417	349	137	60	
4	Alaska Pacific University	Yes	193	146	55	16	
...	
772	Worcester State College	No	2197	1515	543	4	
773	Xavier University	Yes	1959	1805	695	24	
774	Xavier University of Louisiana	Yes	2097	1915	695	34	
775	Yale University	Yes	10705	2453	1317	95	
776	York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	\
0	52	2885	537	7440	3300	450	
1	29	2683	1227	12280	6450	750	
2	50	1036	99	11250	3750	400	
3	89	510	63	12960	5450	450	
4	44	249	869	7560	4120	800	
...	
772	26	3089	2029	6797	3900	500	
773	47	2849	1107	11520	4960	600	
774	61	2793	166	6900	4200	617	
775	99	5217	83	19840	6510	630	
776	63	2988	1726	4990	3560	500	

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend	Grad.Rate
0	2200	70	78	18.1	12	7041	60
1	1500	29	30	12.2	16	10527	56
2	1165	53	66	12.9	30	8735	54
3	875	92	97	7.7	37	19016	59
4	1500	76	72	11.9	2	10922	15
...
772	1200	60	60	21.0	14	4469	40
773	1250	73	75	13.3	31	9189	83
774	781	67	75	14.4	20	8323	49
775	2115	96	96	5.8	49	40386	99


```
776      1250    75      75      18.1      28    4509      99
```

```
[777 rows x 19 columns]
```

You should notice that the first column is just the name of each university in a column named something like `Unnamed: 0`. We don't really want pandas to treat this as data. However, it may be handy to have these names for later. Try the following commands and similarly look at the resulting data frames.

```
[10]: college2 = pd.read_csv('../data/College.csv', index_col=0)
      college3 = college2.rename({'Unnamed: 0': 'College'}, axis=1)
      college3 = college3.set_index('College')
      college2
```

```
[10]:
```

	Private	Apps	Accept	Enroll	Top10perc	\
Abilene Christian University	Yes	1660	1232	721	23	
Adelphi University	Yes	2186	1924	512	16	
Adrian College	Yes	1428	1097	336	22	
Agnes Scott College	Yes	417	349	137	60	
Alaska Pacific University	Yes	193	146	55	16	
...	
Worcester State College	No	2197	1515	543	4	
Xavier University	Yes	1959	1805	695	24	
Xavier University of Louisiana	Yes	2097	1915	695	34	
Yale University	Yes	10705	2453	1317	95	
York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	\
Abilene Christian University	52	2885	537	7440	
Adelphi University	29	2683	1227	12280	
Adrian College	50	1036	99	11250	
Agnes Scott College	89	510	63	12960	
Alaska Pacific University	44	249	869	7560	
...	
Worcester State College	26	3089	2029	6797	
Xavier University	47	2849	1107	11520	
Xavier University of Louisiana	61	2793	166	6900	
Yale University	99	5217	83	19840	
York College of Pennsylvania	63	2988	1726	4990	

	Room.Board	Books	Personal	PhD	Terminal	\
Abilene Christian University	3300	450	2200	70	78	
Adelphi University	6450	750	1500	29	30	
Adrian College	3750	400	1165	53	66	
Agnes Scott College	5450	450	875	92	97	
Alaska Pacific University	4120	800	1500	76	72	
...	
Worcester State College	3900	500	1200	60	60	

Xavier University	4960	600	1250	73	75
Xavier University of Louisiana	4200	617	781	67	75
Yale University	6510	630	2115	96	96
York College of Pennsylvania	3560	500	1250	75	75

	S.F.Ratio	perc.alumni	Expend	Grad.Rate
Abilene Christian University	18.1	12	7041	60
Adelphi University	12.2	16	10527	56
Adrian College	12.9	30	8735	54
Agnes Scott College	7.7	37	19016	59
Alaska Pacific University	11.9	2	10922	15
...
Worcester State College	21.0	14	4469	40
Xavier University	13.3	31	9189	83
Xavier University of Louisiana	14.4	20	8323	49
Yale University	5.8	49	40386	99
York College of Pennsylvania	18.1	28	4509	99

[777 rows x 18 columns]

```
[11]: college3
```

```
[11]:
```

	Private	Apps	Accept	Enroll	Top10perc	\
College						
Abilene Christian University	Yes	1660	1232	721	23	
Adelphi University	Yes	2186	1924	512	16	
Adrian College	Yes	1428	1097	336	22	
Agnes Scott College	Yes	417	349	137	60	
Alaska Pacific University	Yes	193	146	55	16	
...	
Worcester State College	No	2197	1515	543	4	
Xavier University	Yes	1959	1805	695	24	
Xavier University of Louisiana	Yes	2097	1915	695	34	
Yale University	Yes	10705	2453	1317	95	
York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	\
College					
Abilene Christian University	52	2885	537	7440	
Adelphi University	29	2683	1227	12280	
Adrian College	50	1036	99	11250	
Agnes Scott College	89	510	63	12960	
Alaska Pacific University	44	249	869	7560	
...	
Worcester State College	26	3089	2029	6797	
Xavier University	47	2849	1107	11520	
Xavier University of Louisiana	61	2793	166	6900	

Yale University	99	5217	83	19840
York College of Pennsylvania	63	2988	1726	4990

	Room.Board	Books	Personal	PhD	Terminal	\
College						
Abilene Christian University	3300	450	2200	70	78	
Adelphi University	6450	750	1500	29	30	
Adrian College	3750	400	1165	53	66	
Agnes Scott College	5450	450	875	92	97	
Alaska Pacific University	4120	800	1500	76	72	
...	
Worcester State College	3900	500	1200	60	60	
Xavier University	4960	600	1250	73	75	
Xavier University of Louisiana	4200	617	781	67	75	
Yale University	6510	630	2115	96	96	
York College of Pennsylvania	3560	500	1250	75	75	

	S.F.Ratio	perc.alumni	Expend	Grad.Rate
College				
Abilene Christian University	18.1	12	7041	60
Adelphi University	12.2	16	10527	56
Adrian College	12.9	30	8735	54
Agnes Scott College	7.7	37	19016	59
Alaska Pacific University	11.9	2	10922	15
...
Worcester State College	21.0	14	4469	40
Xavier University	13.3	31	9189	83
Xavier University of Louisiana	14.4	20	8323	49
Yale University	5.8	49	40386	99
York College of Pennsylvania	18.1	28	4509	99

[777 rows x 18 columns]

This has used the first column in the file as an index for the data frame. This means that pandas has given each row a name corresponding to the appropriate university. Now you should see that the first data column is **Private**. Note that the names of the colleges appear on the left of the table. We also introduced a new Python object above: a dictionary, which is specified by (key, value) pairs. Keep your modified version of the data with the following:

```
[12]: college = college3
```

(c) Use the `describe()` method of to produce a numerical summary of the variables in the data set.

```
[13]: college.describe()
```

```
[13]:
```

	Apps	Accept	Enroll	Top10perc	Top25perc	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	

mean	3001.638353	2018.804376	779.972973	27.558559	55.796654
std	3870.201484	2451.113971	929.176190	17.640364	19.804778
min	81.000000	72.000000	35.000000	1.000000	9.000000
25%	776.000000	604.000000	242.000000	15.000000	41.000000
50%	1558.000000	1110.000000	434.000000	23.000000	54.000000
75%	3624.000000	2424.000000	902.000000	35.000000	69.000000
max	48094.000000	26330.000000	6392.000000	96.000000	100.000000

	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books \
count	777.000000	777.000000	777.000000	777.000000	777.000000
mean	3699.907336	855.298584	10440.669241	4357.526384	549.380952
std	4850.420531	1522.431887	4023.016484	1096.696416	165.105360
min	139.000000	1.000000	2340.000000	1780.000000	96.000000
25%	992.000000	95.000000	7320.000000	3597.000000	470.000000
50%	1707.000000	353.000000	9990.000000	4200.000000	500.000000
75%	4005.000000	967.000000	12925.000000	5050.000000	600.000000
max	31643.000000	21836.000000	21700.000000	8124.000000	2340.000000

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni \
count	777.000000	777.000000	777.000000	777.000000	777.000000
mean	1340.642214	72.660232	79.702703	14.089704	22.743887
std	677.071454	16.328155	14.722359	3.958349	12.391801
min	250.000000	8.000000	24.000000	2.500000	0.000000
25%	850.000000	62.000000	71.000000	11.500000	13.000000
50%	1200.000000	75.000000	82.000000	13.600000	21.000000
75%	1700.000000	85.000000	92.000000	16.500000	31.000000
max	6800.000000	103.000000	100.000000	39.800000	64.000000

	Expend	Grad.Rate
count	777.000000	777.000000
mean	9660.171171	65.46332
std	5221.768440	17.17771
min	3186.000000	10.00000
25%	6751.000000	53.00000
50%	8377.000000	65.00000
75%	10830.000000	78.00000
max	56233.000000	118.00000

(d) Use the `pd.plotting.scatter_matrix()` function to produce a scatterplot matrix of the first columns `[Top10perc, Apps, Enroll]`. Recall that you can reference a list `C` of columns of a data frame `A` using `A[C]`.

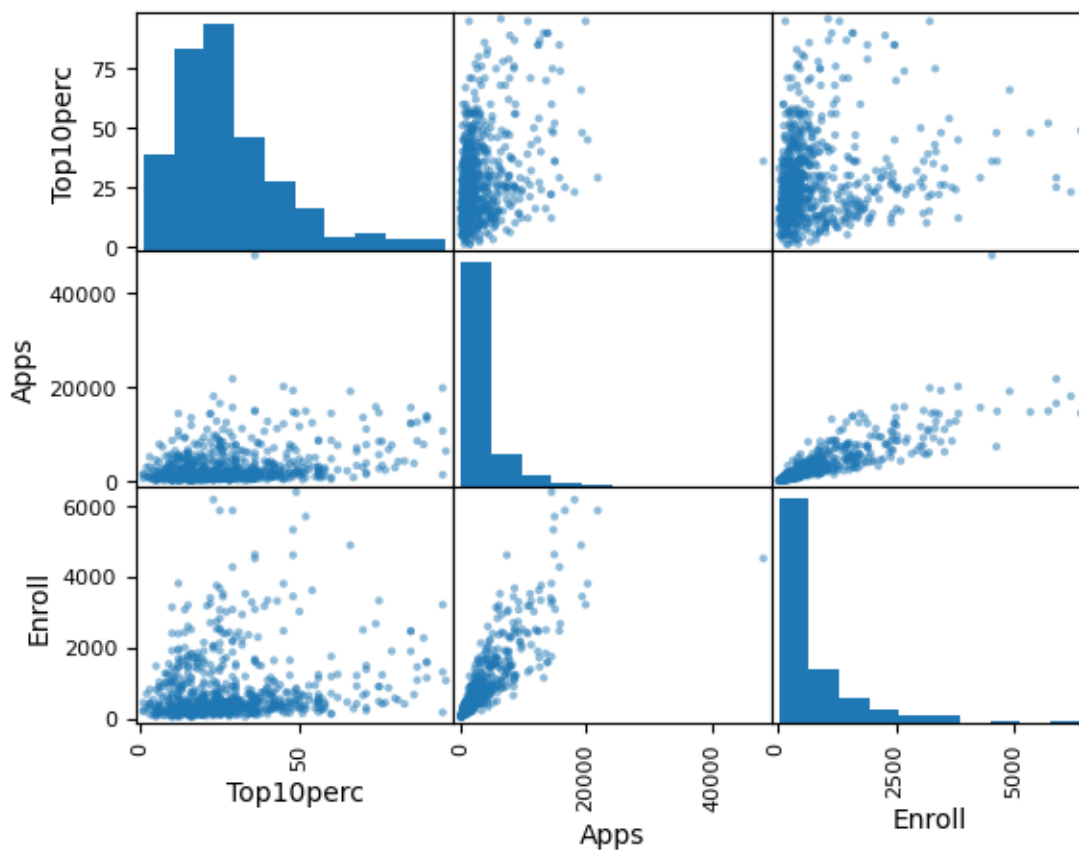
```
[14]: pd.plotting.scatter_matrix(college[['Top10perc', 'Apps', 'Enroll']])
```

```
[14]: array([[<Axes: xlabel='Top10perc', ylabel='Top10perc'>,
             <Axes: xlabel='Apps', ylabel='Top10perc'>,
             <Axes: xlabel='Enroll', ylabel='Top10perc'>],
            [<Axes: xlabel='Top10perc', ylabel='Apps'>,
             <Axes: xlabel='Top10perc', ylabel='Enroll'>],
            [<Axes: xlabel='Apps', ylabel='Apps'>,
             <Axes: xlabel='Apps', ylabel='Enroll'>],
            [<Axes: xlabel='Enroll', ylabel='Enroll'>]])
```

```

<Axes: xlabel='Apps', ylabel='Apps'>,
<Axes: xlabel='Enroll', ylabel='Apps'>],
[<Axes: xlabel='Top10perc', ylabel='Enroll'>,
<Axes: xlabel='Apps', ylabel='Enroll'>,
<Axes: xlabel='Enroll', ylabel='Enroll'>]], dtype=object)


```



(e) Use the `boxplot()` method of `college` to produce side-by-side boxplots of `Outstate` versus `Private`.

```
[15]: college.boxplot(column='Outstate', by='Private')
```

```
[15]: <Axes: title={'center': 'Outstate'}, xlabel='Private'>
```



(f) Create a new qualitative variable, called Elite, by binning the Top10perc variable into two groups based on whether or not the proportion of students coming from the top 10% of their high school classes exceeds 50%.

```
[16]: college['Elite'] = pd.cut(college['Top10perc'],
                                [0,50,100],
                                labels=['No', 'Yes'])
```

Use the value_counts() method of college['Elite'] to see how many elite universities there are.


```
[17]: college['Elite'].value_counts()
```

```
[17]: Elite
No      699
Yes      78
Name: count, dtype: int64
```

Finally, use the boxplot() method again to produce side-by-side boxplots of Outstate versus Elite.

```
[18]: college.boxplot(column='Outstate', by='Elite')
```

```
[18]: <Axes: title={'center': 'Outstate'}, xlabel='Elite'>
```

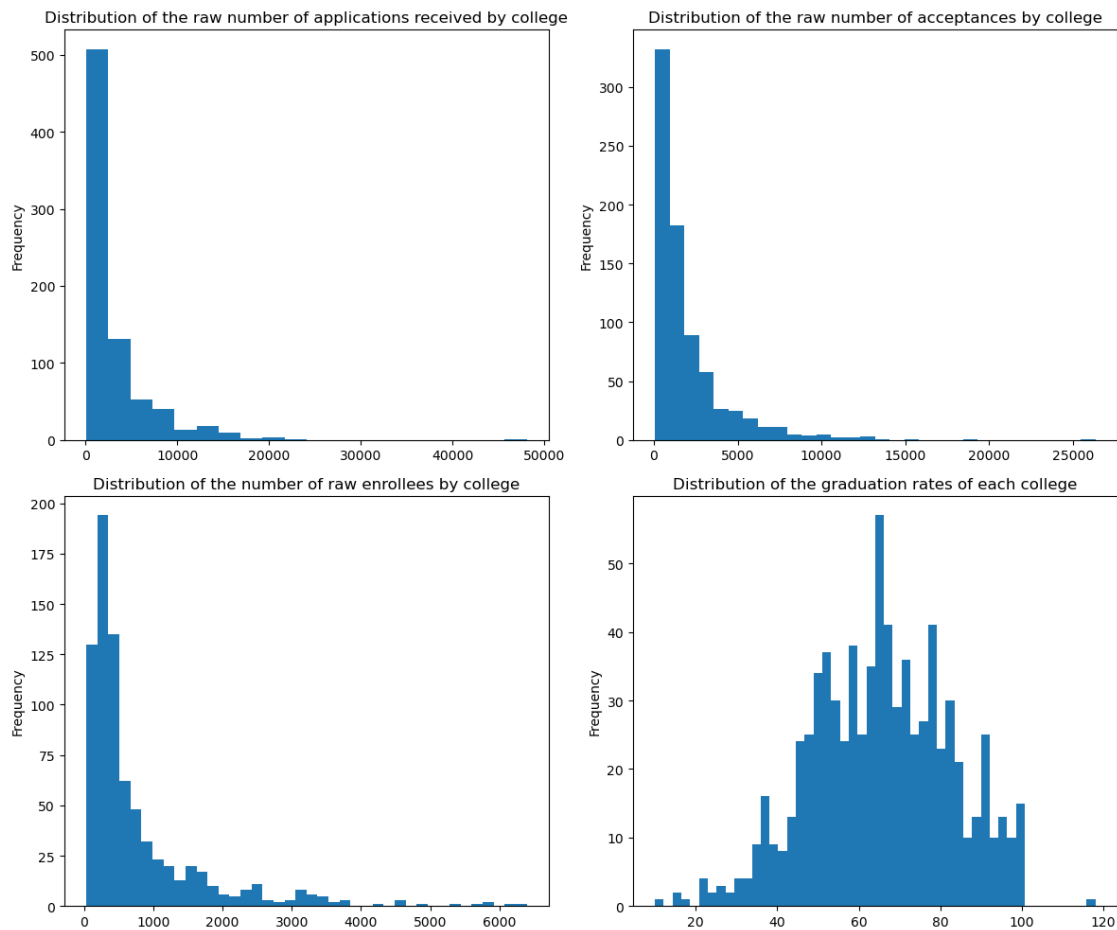


(g) Use the `plot.hist()` method of `college` to produce some histograms with differing numbers of bins for a few of the quantitative variables. The command `plt.subplots(2, 2)` may be useful: it will divide the plot window into four regions so that four plots can be made simultaneously. By changing the arguments you can divide the screen up in other combinations.

```
[19]: fig, axes = plt.subplots(2, 2, figsize=(12, 10))
      ax1, ax2, ax3, ax4 = axes.flatten()

      college['Apps'].plot.hist(ax=ax1, bins=20, title='Distribution of the raw number_
      ↳of applications received by college')
      college['Accept'].plot.hist(ax=ax2, bins=30, title='Distribution of the raw_
      ↳number of acceptances by college')
      college['Enroll'].plot.hist(ax=ax3, bins=40, title='Distribution of the number_
      ↳of raw enrollees by college')
      college['Grad.Rate'].plot.hist(ax=ax4, bins=50, title='Distribution of the_
      ↳graduation rates of each college')
```

```
plt.tight_layout()
```




(h) Continue exploring the data and provide a brief summary of what you discover.

```
[20]: select = ['Outstate', 'F.Undergrad', 'Enroll', 'Grad.Rate', 'PhD', 'S.F.Ratio', '
        ↪ 'perc.alumni', 'Expend']
pd.plotting.scatter_matrix(college[select], figsize=(12, 12))
plt.suptitle('Pair Plot of College Data')
```

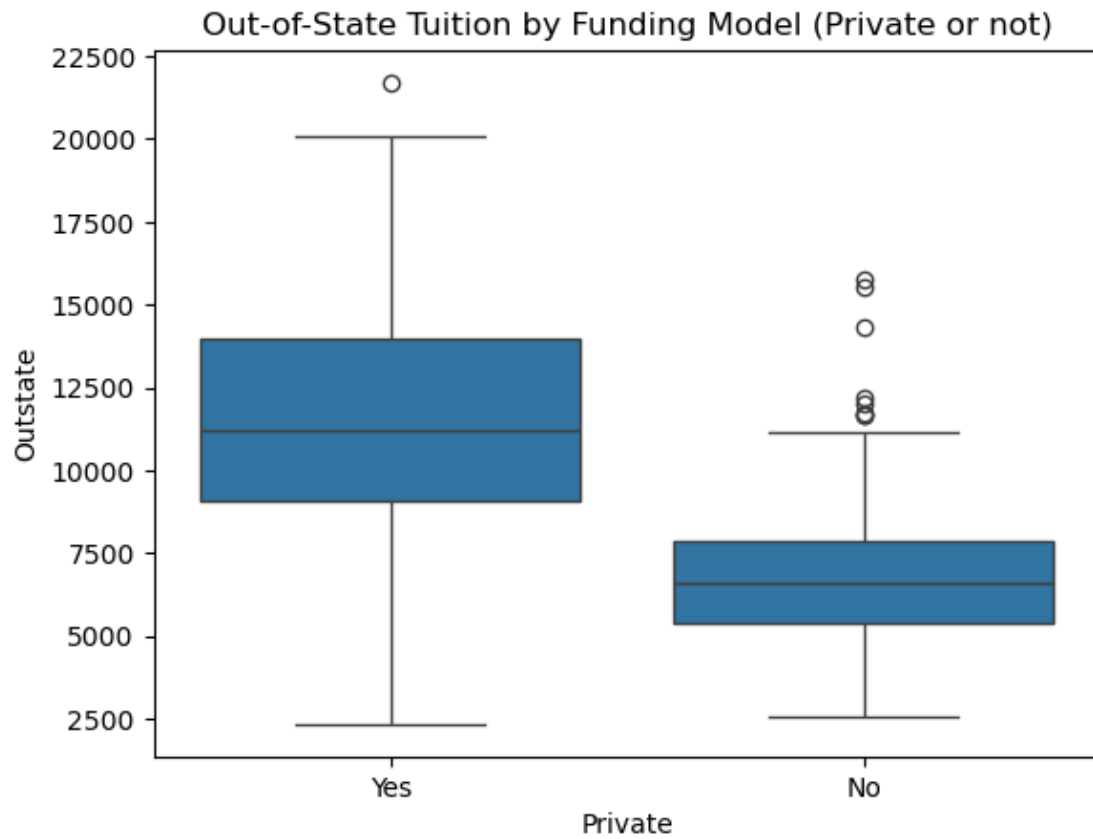
```
[20]: Text(0.5, 0.98, 'Pair Plot of College Data')
```


Pair Plot of College Data



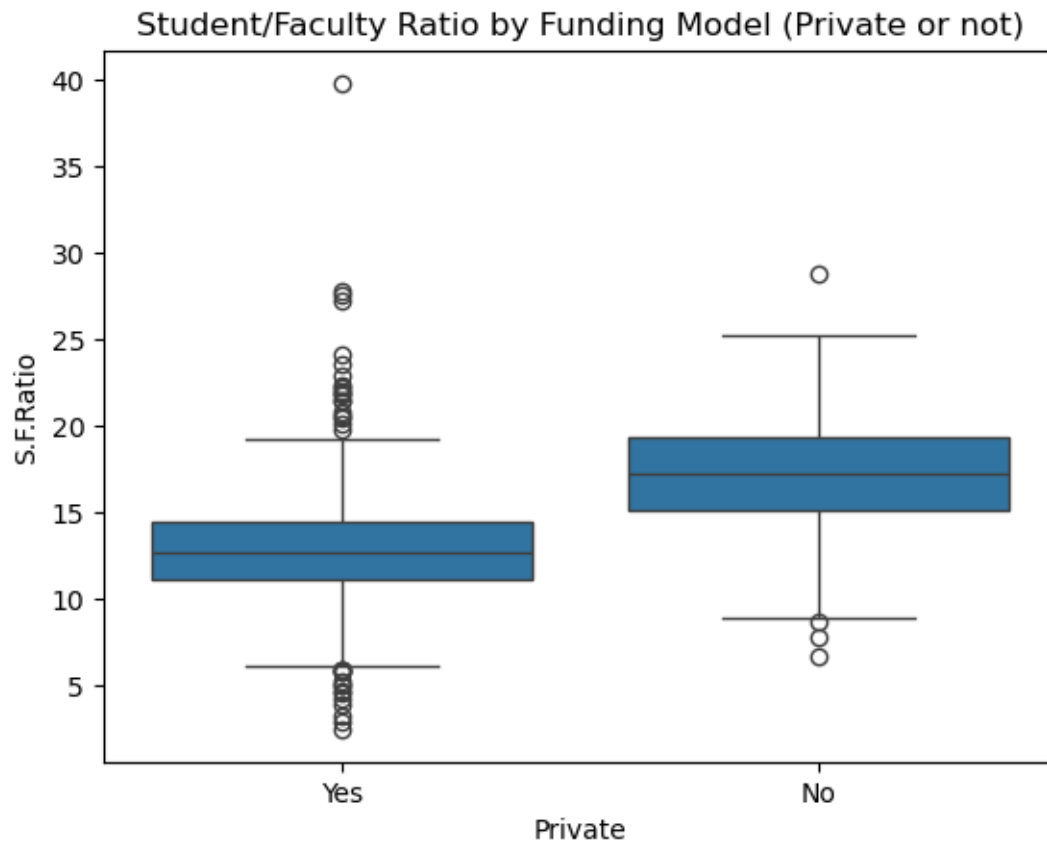
```
[21]: sns.boxplot(data=college, x='Private', y='Outstate')
plt.title('Out-of-State Tuition by Funding Model (Private or not)')
```

```
[21]: Text(0.5, 1.0, 'Out-of-State Tuition by Funding Model (Private or not)')
```



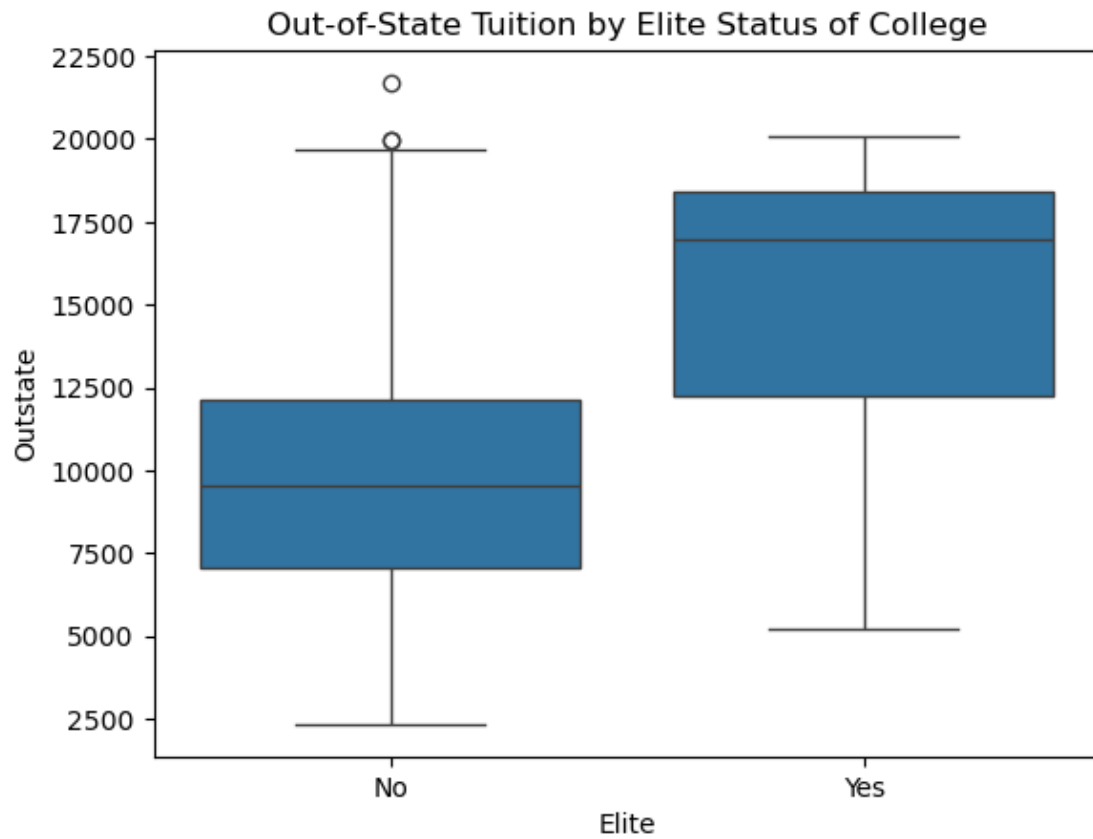
```
[22]: sns.boxplot(data=college, x='Private', y='S.F.Ratio')  
plt.title('Student/Faculty Ratio by Funding Model (Private or not)')
```

```
[22]: Text(0.5, 1.0, 'Student/Faculty Ratio by Funding Model (Private or not)')
```




```
[23]: sns.boxplot(data=college, x='Elite', y='Outstate')
plt.title('Out-of-State Tuition by Elite Status of College')
```

```
[23]: Text(0.5, 1.0, 'Out-of-State Tuition by Elite Status of College')
```



```
[24]: sns.boxplot(x='Elite', y='Grad.Rate', data=college)
plt.title('Graduation Rate by Elite Status of College')
```

```
[24]: Text(0.5, 1.0, 'Graduation Rate by Elite Status of College')
```



The pair plot matrix provides a visual representation of the relationships between several key quantitative variables. There is a strong positive correlation between per-student instructional expenditure and the out-of-state tuition of colleges. Generally, universities with higher out-of-state tuition typically spend more money to teach the average student at their institution. While there is an almost linear relationship between total full-time undergraduate enrollment and enrollment for the year, this could be attributed to multicollinearity since most colleges typically have some relationship between students enrolled in a year and the total undergraduate class size. I can also observe a strong negative correlation between student/faculty ratio and out-of-state tuition, indicating that schools with fewer students per faculty member charge higher tuition for out-of-state students.

Private schools tend to have higher out-of-state tuition and lower student/faculty ratios. On the other hand “elite” colleges tend to have a higher graduation rate, and a higher out-of-state tuition.